

# PROJECT : IMAGE DENOISING

By: Harsh Pratap Rathore, EE 3Y (22115063)

## **Introduction**

This project focuses on utilizing deep learning techniques to enhance low-light images, improving their clarity and usability. The project involves training a model using a dataset of low-light and high-light images, evaluating its performance using metrics such as PSNR (Peak Signal-to-Noise Ratio), and deploying the model for inference. The following report details the steps taken, the methodology employed.

## **Mounting the drive containing dataset**

The initial step in our project involves importing the dataset, which is stored on Google Drive. We need to mount the Google Drive to access the dataset files.

## **Importing Necessary Libraries**

- `os`, `random`, `numpy`: For file handling, randomization, and numerical operations.
- `glob`: To retrieve files matching specified patterns.
- `PIL` (Python Imaging Library), `ImageOps`: For image manipulation.
- `matplotlib.pyplot`: For plotting images.
- `cv2` (OpenCV): For image processing tasks.
- `tensorflow`, `keras`, `layers`: For building and training neural network models.

## **Image Preprocessing Function**

This function is responsible for loading and preprocessing images.

- **`tf.io.read_file(image_path)`**: This function reads the image file from the specified path as a byte string.
- **`tf.image.decode_png(image, channels=3)`**: Decodes the byte string to a PNG image tensor. The “channels=3” argument ensures the image is loaded with RGB.
- **`tf.image.resize(images=image, size=[256, 256])`**: Resizes the image to 256x256 pixels, a standard size for input into neural networks.
- **`image / 255.0`**: Normalizes the pixel values of the image to the range [0, 1].

## **Data Pipeline Creation Function**

This function efficiently loads, preprocess, and batch the images for training or inference.

- **`tf.data.Dataset.from_tensor_slices(low_light_images)`**: This function creates a “`tf.data.Dataset`” from the list of image paths for low\_light images.
- **`dataset.map(preprocess_image, num_parallel_calls=tf.data.AUTOTUNE)`**: The map function applies the “preprocess\_image” function to each element in the dataset. The “num\_parallel\_calls=tf.data.AUTOTUNE” argument enables parallel processing, optimizing the pipeline's performance by utilizing multiple CPU cores.
- **`dataset.batch(16, drop_remainder=True)`**: Batches the dataset into groups of 16 images.

## **Dataset Preparation**

- **Image Path Retrieval:** We will take first 30 images for testing our results, from 31 to 400 images for training and from 400 onwards for validation.
- **Dataset Creation:**
  - `create_dataset(train_low_light_images)`: Preprocesses and batches training images.
  - `create_dataset(val_low_light_images)`: Preprocesses and batches validation images.
- **Print Dataset Details:** Outputs information about the datasets to confirm correct loading and preprocessing.

## **Building the Deep Curve Estimation (DCE) Network**

This function “`build_image_enhancement_net()`” defines a Deep Curve Estimation network architecture using the Keras API from TensorFlow.

- **Input Layer:** Defines the “`input_image`” variable for the network with “`shape [None, None, 3]`”.
- **Convolutional Layers:** Sequentially defines several convolutional layers (`conv_layer1` to `conv_layer4`). Each layer uses a 3x3 kernel, “`relu`” activation function, and same padding to ensure that the output has the same spatial dimensions as the input.
- **Intermediate Concatenation Layers:** Concatenates the output of one with another along the last axis (-1), corresponding to the channel axis.
- **Output Layer:** Final convolutional layer with 24 filters, 3x3 kernel size, `tanh` activation function, and same padding.
- **Model Compilation:** Returns a Keras Model object initialized with the input and output layers. This model defines the complete architecture of the DCE network.

## **Function to Calculate Color Constancy Loss**

- **Calculate Mean RGB Values:** Computes the mean red (`mean_r`), green (`mean_g`), and blue (`mean_b`) values across the image batch using `tf.reduce_mean`.
- **Calculate Squared Differences:** Computes `d_rg`, `d_rb`, and `d_gb` as squared differences between mean RGB values.
- **Calculate Color Constancy Loss:** Computes the color constancy loss as  $\sqrt{d_{rg}^2 + d_{rb}^2 + d_{gb}^2}$ , aiming to maintain color consistency across images.

## **Function to Calculate Exposure Loss**

- **Calculate Mean Intensity:** Computes the mean intensity of the input `image_tensor` across the RGB channels using `tf.reduce_mean`.
- **Average Pooling:** Performs average pooling on `mean_intensity` to compute the overall mean with a kernel size (`ksize`) of 16x16 and strides (`strides`) of 16.
- **Calculate Exposure Loss:** Computes the exposure loss as the mean squared difference between `pooled_mean` and `target_mean`. This loss helps in adjusting the exposure of the images.

### **Function to Calculate Illumination Smoothness Loss**

- **Batch Size and Image Dimensions:** Retrieves the batch size (batch\_size), height (h\_x), and width (w\_x) of the image\_tensor.
- **Number of Horizontal and Vertical Pairwise Differences:** Computes count\_h and count\_w as the total number of horizontal and vertical pairwise differences respectively.
- **Calculate Horizontal and Vertical Total Variation (TV):** Computes h\_tv and w\_tv as the sum of squared differences between horizontally and vertically adjacent pixels across all channels in each image of the batch respectively.
- **Convert Counts to Float32:** Converts batch\_size, count\_h, and count\_w to float32 for division to ensure compatibility with TensorFlow operations.
- **Calculate Illumination Smoothness Loss:**

$$\text{illumination\_smoothness\_loss} = \frac{2 \times (h\_tv / \text{count\_h} + w\_tv / \text{count\_w})}{\text{batch\_size}}$$

- This loss measures how smoothly illumination changes across adjacent pixels in the image tensor, promoting spatial consistency in illumination.

### **Spatial Consistency Loss Class**

- **Constructor (\_\_init\_\_ method):**
  - Initializes the class as a subclass of keras.losses.Loss with reduction="none", indicating that the loss will be computed per batch sample without reducing across the batch.
  - Defines four convolutional kernels (left\_kernel, right\_kernel, up\_kernel, down\_kernel) using TensorFlow constants (tf.constant). These kernels are used to compute differences in the spatial consistency of illumination between the original (y\_true) and enhanced (y\_pred) images.
- **Mean Intensity Calculation:** Computes the mean intensity (original\_mean, enhanced\_mean) of the original (y\_true) and enhanced (y\_pred) images along the channel axis (axis=3) using tf.reduce\_mean.
- **Average Pooling:** Performs average pooling (tf.nn.avg\_pool2d) on original\_mean and enhanced\_mean with a kernel size (ksize) of 4x4 and strides of 4x4 (strides=[1, 4, 4, 1]) to obtain pooled representations (original\_pool, enhanced\_pool) with reduced spatial dimensions.
- **Spatial Differences Computation:**
  - Computes spatial differences (d\_original\_left, d\_original\_right, d\_original\_up, d\_original\_down) between adjacent pixels in original\_pool using convolution operations (tf.nn.conv2d) with predefined kernels (left\_kernel, right\_kernel, up\_kernel, down\_kernel).
  - Similarly, computes spatial differences (d\_enhanced\_left, d\_enhanced\_right, d\_enhanced\_up, d\_enhanced\_down) between adjacent pixels in enhanced\_pool.
- **Squared Differences:** Computes the squared differences (d\_left, d\_right, d\_up, d\_down) between d\_original\_\* and d\_enhanced\_\* tensors.

It returns the sum of squared differences ( $d_{\text{left}} + d_{\text{right}} + d_{\text{up}} + d_{\text{down}}$ ) and measures how well the enhanced image preserves the spatial arrangement of illumination from the original image.

### **Zero-DCE Model Class**

- **Constructor(`__init__` method):** Inherits from `keras.Model` and initializes the `dce_model` attribute by calling `build_image_enhancement_net()`, which constructs the Deep Curve Estimation (DCE) network for image enhancement.
- **Compilation (compile method):**
  - Overrides the `compile` method of `keras.Model`.
  - Configures the optimizer as Adam with a specified `learning_rate`.
  - Initializes `spatial_constancy_loss` as an instance of `SpatialConsistencyLoss` with `reduction="none"`.
- **Image Enhancement (get\_enhanced\_image method):** Enhances the input data using the outputs (output) from the DCE network.
  - Splits output into eight components (`r1` to `r8`).
  - Iteratively enhances data using these components based on a formula involving element-wise operations with the components.
- **Forward Pass (call) Method:**
  - Overrides the `call` method of `keras.Model`.
  - Executes the DCE network (`dce_model`) on the input data and enhances the image using `get_enhanced_image`.
- **Training Step (train\_step method):**
  - Overrides the `train_step` method of `keras.Model`.
  - Computes gradients of the total loss with respect to trainable weights of `dce_model` using `tf.GradientTape`.
  - Updates weights of `dce_model` using the optimizer based on computed gradients.
  - Returns losses computed during the step.
- **Testing (test\_step method):**
  - Overrides the `test_step` method of `keras.Model`.
  - Performs forward pass through `dce_model` to get outputs and computes losses using `compute_losses`.
  - Returns losses computed during testing.

### **Training the Zero-DCE Model**

- **Model Initialization:** It creates an instance of the Zero-DCE model.
- **Compilation:**
  - `zero_dce_model.compile(learning_rate=1e-4)`: Configures the model for training.
  - Uses Adam optimizer with a learning rate of `1e-4`.
  - Initializes losses including `spatial_constancy_loss` defined in `SpatialConsistencyLoss` class.

- **Training:** It trains the model on train\_dataset and evaluates on val\_dataset for 70 epochs and returns a history object that contains training/validation loss values over epochs.

### **“perform\_inference” Function**

- Converts the PIL image to a numpy array (image\_array).
- Normalizes pixel values to the range [0, 1].
- Expands the dimensions to match the model’s input shape (adds a batch dimension).
- Passes the preprocessed image through the Zero-DCE model (zero\_dce\_model).
- Scales the output image back to the [0, 255] range and converts it to uint8.
- Converts the numpy array back to a PIL image (output\_image).
- output\_image: Enhanced output image as a PIL image.

### **Summary of Findings**

In this project, we implemented the Zero-Reference Deep Curve Estimation (Zero-DCE) network for enhancing low-light images. The Zero-DCE model estimates light enhancement curves (LE-curves) to improve image quality iteratively. Key outcomes include:

- **Enhanced Image Quality:** The model effectively enhances brightness and contrast while maintaining natural colors and details.
- **Effective Loss Functions:** Utilized non-reference loss functions such as spatial consistency, exposure control, color constancy, and illumination smoothness to guide the enhancement process.
- **Quantitative Improvement:** The model's performance, measured by PSNR, showed significant improvement in image quality.

### **Methods to Further Improve the Project**

1. **Enhanced Network Architecture:** Experiment with deeper networks and integrate attention mechanisms.
2. **Advanced Loss Functions:** Incorporate perceptual and adversarial loss functions for more visually pleasing and realistic enhancements.
3. **Data Augmentation:** Use larger, more diverse datasets and advanced augmentation techniques to improve model robustness.
4. **Real-time Enhancements:** Optimize for faster inference times and develop lightweight versions for mobile and embedded devices.
5. **User Interactivity:** Allow users to adjust enhancement parameters and gather feedback to refine the model.
6. **Cross-domain Applications:** Extend the model for video enhancement and multispectral imaging applications.

The link of the research paper I took help from is given

<https://ar5iv.labs.arxiv.org/html/2001.06826>

We have tested first 30 images and displayed the corresponding predicted and high light images along with their psnr value. **The average psnr value in my results is around 28.21 dB.**